

Question No 02:Slide 25:Solution 1: $\pi_{sname} ((G_{bid=103} Reserves) \bowtie Sailors)$ → first we solve $(G_{bid=103} Reserves)$

<u>Sid</u>	<u>bid</u>	<u>day</u>
58	103	11/12/96

Now

 $((G_{bid=103} Reserves) \bowtie Sailors)$

<u>Sid</u>	<u>Sname</u>	<u>rating</u>	<u>age</u>	<u>bid</u>	<u>day</u>
58	Rusty	10	35	103	11/12/96

Now

 $\pi_{sname} ((G_{bid=103} Reserves) \bowtie Sailors)$

<u>Sname</u>
Rusty

— xp — xp — xp — xp

— xp — xp

Slide 26:

$\pi_{\text{sname}} ((\sigma_{\text{color} = 'red'} \text{Boats}) \bowtie \text{Reserves})$
 $\bowtie \text{Sailors}$

→ first we solve $(\sigma_{\text{color} = 'red'} \text{Boats})$

bid	bname	color
102	Interlake	Red
104	Marine	Red

Now

$(\sigma_{\text{color} = 'red'} \text{Boats}) \bowtie \text{Reserves}$

bid	bname	color	sid	day

Now

$(\sigma_{\text{color} = 'red'} \text{Boats}) \bowtie \text{Reserves} \bowtie \text{Sailors}$

sid	sname	rating	age	bid	day	bname	color

Now

$\pi_{\text{sname}} ((\sigma_{\text{color} = 'red'} \text{Boats}) \bowtie \text{Reserves} \bowtie \text{Sailors})$

sname

— No — No — No — No

— No

3

Date: __/__/20__

Method 2

$\pi_{\text{Sname}} (\pi_{\text{Sid}} ((\pi_{\text{bid}} \text{ Gcolor} = \text{'red' Boats}) \bowtie \text{Res}) \bowtie \text{Sailers}))$

→ first we solve $(\pi_{\text{bid}} \text{ Gcolor} = \text{'red' Boats})$

bid
102
104

Now $(\pi_{\text{bid}} \text{ Gcolor} = \text{'red' Boats}) \bowtie \text{Reserves}$

bid	Sid	day

~~Now $(\pi_{\text{bid}} \text{ Gcolor} = \text{'red' Boats}) \bowtie \text{Reserves}$~~

Now $\pi_{\text{Sid}} ((\pi_{\text{bid}} \text{ Gcolor} = \text{'red' Boats}) \bowtie \text{Reserve})$

Sid

Now

$(\pi_{\text{Sid}} ((\pi_{\text{bid}} \text{ Gcolor} = \text{'red' Boats}) \bowtie \text{Reserve}) \bowtie \text{Sailers}))$

Sid	Sname	rating	age

Now

$\pi_{\text{Sname}} (\pi_{\text{Sid}} ((\pi_{\text{bid}} \text{ Gcolor} = \text{'red' Boats}) \bowtie \text{Reserve}) \bowtie \text{Sailers}))$

Sname

— xp — xp — xp — xp — xp

— xp — xp — xp — xp

(4)

Slide-27:

$\rho(\text{Tempboats}, (\sigma_{\text{color} = 'red' \vee \text{color} = 'green'} \text{Boats}))$
 $\pi_{\text{sname}}(\text{Tempboats} \bowtie \text{Reserves} \bowtie \text{Sailors})$
→ first we solve $(\sigma_{\text{color} = 'red' \vee \text{color} = 'green'} \text{Boats})$

bid	bname	color
102	Interlake	Red
103	Clipper	Green
104	Marine	Red

Now $\rho(\text{Tempboats}, (\sigma_{\text{color} = 'red' \vee \text{color} = 'green'} \text{Boats}))$

Tempboats

bid	bname	color
102	Interlake	Red
103	Clipper	Green
104	Marine	Red

Now

Tempboats \bowtie Reserves

bid	sid	day	bname	color
103	58	11/12/96	Clipper	Green

Now

Tempboats \bowtie Reserves \bowtie Sailors

sid	sname	rating	age	bid	day	bname	color
58	Rusty	10	35.0	103	11/12/96	clipper	Green

Now $\pi_{\text{sname}}(\text{TAR} \bowtie \text{S})$

sname
Rusty

— xx — xx — xx — xx — xx

(5)

Slide - 28:

- $\sigma(\text{Tempred}, \pi_{\text{sid}}((G_{\text{color}} = \text{'red' Boats}) \bowtie \text{Reserves}))$
- $\sigma(\text{Tempgreen}, \pi_{\text{sid}}((G_{\text{color}} = \text{'green' Boats}) \bowtie \text{Reserves}))$
- $\pi_{\text{sname}}((\text{Tempred} \cap \text{Tempgreen}) \bowtie \text{Sailors})$

Sol

First we solve $(G_{\text{color}} = \text{'red' Boats})$

<u>bid</u>	bname	color
102	Interlake	Red
104	Marine	Red

Now

$((G_{\text{color}} = \text{'red' Boats}) \bowtie \text{Reserves})$

<u>bid</u>	<u>Sid</u>	<u>day</u>	bname	color

Now

$\sigma(\text{Tempred}, \pi_{\text{sid}}((\text{ — }) \bowtie R))$

Tempred

<u>Sid</u>

Now

$(G_{\text{color}} = \text{'green' Boats})$

<u>bid</u>	bname	color
103	Clipper	Green

P e r p o

||

(6)

Now

$((G_{color} = 'green', Boats) \bowtie Reserves)$

bld	Sid	day	bname	color
103	58	11/12/96	clipper	green

Now

$\sigma_{(Tempgreen, \wedge sid)}((G_{color} = 'green', Boats) \bowtie Reserves)$

Tempgreen

Sid
58

Now

$(Tempred \wedge Tempgreen)$

Sid
58

Now

$(Tempred \wedge Tempgreen) \bowtie Sailors$

Sid	Sname	rating	age
58	rusty	10	35.0

Now

$\wedge sname ((Tempred \wedge Tempgreen) \bowtie Sailors)$

Sname
rusty

— xp — xp — xp — xp — xp — xp

p + 1 p 0

Slide - 29:

$$\textcircled{1} \quad \rho(\text{TempSids}, (\pi_{\text{sid}, \text{bid Reserves}}) / (\pi_{\text{bid Boats}}))$$

$\pi_{\text{name}}(\text{TempSids} \bowtie \text{Sailors})$

→ first we solve $(\pi_{\text{sid}, \text{bid Reserves}})$

Sid	bid
22	101
58	103

Now $(\pi_{\text{bid Boats}})$

bid
101
102
103
104

Now $(\pi_{\text{sid}, \text{bid Reserves}}) / (\pi_{\text{bid Boats}})$

Sid

$$\rho(\text{tempSids}, (\pi_{\text{sid}, \text{bid Reserves}}) / (\pi_{\text{bid Boats}}))$$

tempSids

Sid

$\rho \rightarrow \pi \rightarrow \rho$

(8)

Now $(\pi_{sid} \bowtie Sailors)$

<u>sid</u>	Sname	nativity	age

Now

$\pi_{sname} (\pi_{sid} \bowtie Sailors)$

<u>Sname</u>

— π — π — π — π — π —

(2) $\rho(Tempsids, (\pi_{sid, bid} Reserves))$

$\wedge \pi_{bid} (\sigma_{bname = 'Interlake' Boats})$

$\pi_{sname} (TempSids \bowtie Sailors)$

Sid

$\rightarrow (\pi_{sid, bid} Reserves)$

<u>sid</u>	<u>bid</u>
22	101
58	103

Now

$(\sigma_{bname = 'Interlake' Boats})$

<u>bid</u>	bname	color
101	Interlake	Blue
102	Interlake	Red

Now $\pi_{bid} (\sigma_{bname = 'Interlake' Boats})$

<u>bid</u>
101
102

ρ , ρ , ρ , ρ ,

(9)

Now $(\pi_{sid, bid} Reserve) / \pi_{bid} (G_{bname} = 'Interlake' Boats)$

Sid

Now

$\rho (TempSide, (\pi_{sid, bid} Reserve) / \pi_{bid} (G_{bname} = 'Interlake' Boats))$

TempSide

Sid

Now

TempSide \bowtie Sailors

Sid	Sname	rating	age

— xp — xp — xp — ~~x~~ — bi

bid	bid
101	22
835	22

(Join of 'Sailors' & 'Boats')

sid	sname	rating
101	Interlake	10
835	Interlake	10

(Join of 'Sailors' & 'Boats')

bid
101
835

Question: 03

A) File Organization & Indexing:

(1) Ans: Following are three alternative file organizations.

a) Heap Files: Suitable when typical access is file scan of all records.

b) Sorted Files:

→ Best for retrieval in search key order.

→ also good for search based on search key.

c) Indexes:

Organize data via trees or hashing.

(2) Ans: Index: A database index is a disk-based data structure that improves the speed of data retrieval.

Indexing organize records via trees or hashing.

(3): (1) Selection of form field (op) Constant
 (2) Equality Selections (op is =)
 (3) Range Selections (op is one of <, >, =)
 & <=, BETWEEN)
 (4) One common n-dimensional index: R-tree

(2)


- (4)
- (1) data record with key value K
 - (2) $\langle K, \text{rid of data record of with search key value } K \rangle$
 - (3) $\langle K, \text{list of rids of data records with search key } K \rangle$

(5) * Good for equality selections.
→ Index is collection of buckets.
Bucket = primary page plus zero or more overflow page.


→ Hashing Function h :
- $h(x)$ = bucket in records x belongs.
- h looks at search key fields of x .

* If alternative cost (1) is used, the bucket contains the data records; otherwise, they contain $\langle \text{key}, \text{rid} \rangle$ or $\langle \text{key}, \text{rid-list} \rangle$ pairs.

(6) The B+ tree is a balanced binary search tree. It follows multi-level index format.



- * Leaf pages contains data entries, and are chained (prev and next).
- * Non leaf pages contains index entries and direct location.

Example:

- * Find 28^* ? 29^* ? All $> 15^*$ and $< 30^*$
- * Insert/delete: Find data Entry in leaf, then change it. Need to adjust parent sometimes.
- And change sometimes bubbles up the tree.

(7) I/O Cost operations: the table shows diff

the number of data pages B , Number of records per page, Fanout of B-tree, and Time required for equality search for different file organization methods, for different operations.


To scan all files records, Heap file needs B number of pages, sorted file needs B number of pages, and clustered and unclustered Tree needs $1.5B$ number of pages while Hash index needs $1.25B$ number of pages.

Similar things can be interpreted from the table for Get all in sorted form, Equality search, Range Search, Insert and delete.

(8) clustered index:

If order of data stored is the same as or 'close to' order of index data entries, then called clustered index.

Example:



B.) Query Processing joins and Sorting

1) The dbms must sort the ~~query~~ data to run the query for

2) Different techniques used by DBMS are

- 1) 3-way sort
- 2) Two-way External merge sort
- 3) General external merge sort.
- 4) Internal sort
- 5) Quick Sort
- 6) Heap sort

3) Ans: Scenario: Table to be sorted has B+ tree index on sorting columns.

Idea: Can retrieve records in order by traversing leaf pages.

Case to be considered:

- B+ tree is clustered Good Idea!
- B+ tree is not clustered Bad Idea.

4) To sort a file with N pages using B buffer pages:

- pass 0: Use B buffer pages produce $\lfloor N/B \rfloor$ sorted runs of B pages each.

- pass 1, 2, ... etc. % merge B-1 runs.

5): Four join algorithms:

- 1) Hash-join
- 2) Sort merge join (RMS)
- 3) Simple Nested loop join
- 4) Index Nested loop join

6) Following are three nested loop joins:

1) Block Nested loop join:

Use one page as an input buffer for scanning the inner S , one page as the output buffer and use all remaining pages to hold block of outer R .

→ for each matching tuple r in R -block, s in S -page, add $\langle r, s \rangle$ to result.

(6)

2) Simple Nested loops join:

→ For each tuple in outer relation relation R , we scan the entire inner relation S .

→ Page oriented Nested loops join:
for each page of R , get each page of S , and write out matching pairs of tuples $\langle r, s \rangle$ where r is in R -page and s is in S -page

3) Index Nested loop join:

- If there is an index on the join column of one relation (say S), can make it the inner and exploit the index.

- For each R tuple, cost of probing S index is about 1.2 for hash index, 2.4 for B+ tree.

1) Sort-Merge Join:

Sort R and S on the join column, then scan them to do a merge, and output the result tuple.
 R is scanned once, each S group is scanned once per matching R tuple.

→ Advance scan of R until current R -tuple $>$ current S tuple, then advance scan of S until current S -tuple $>$ current R tuple: do this until current $R =$ current

S tuple.

8) Hash Join: Partition both relations using hash function h .
 R tuples in partition i will only match S tuples in partition i .
 Read in a partition of R , hash it using $h_2 (< > R h!)$. Scan matching partition of S , search for matches.

(C) Query Optimization:

1)

Hash join Method is faster.

2)

Access path ~~means~~ ^{method of retrieving tuples} "the path a Query ~~can have towards the source table.~~"

3)

a tree index matches terms that involve only attributes in a prefix of the search.

4).

A hash index matches terms that has a term attribute = value for every attribute in the search key of the index.

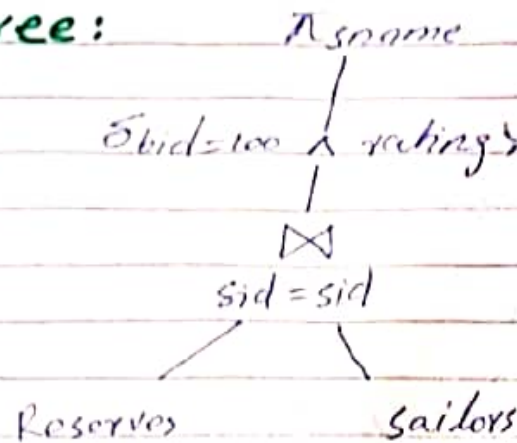
5)

A B+ tree index on day can be used; then, $bid=5$ and $sid=3$ must be checked for each retrieved tuple.

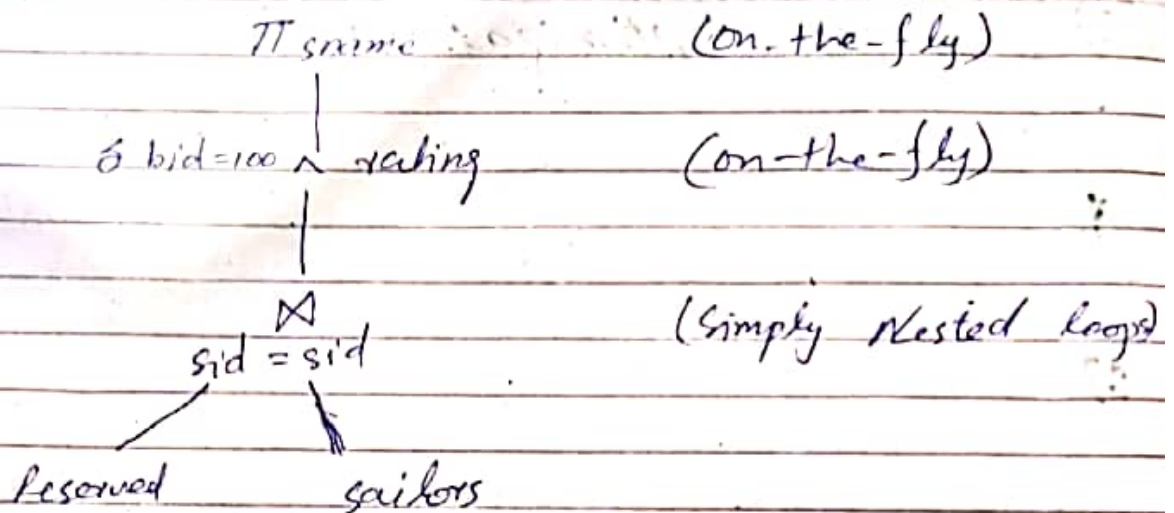
A hash index on $\langle bid, sid \rangle$ could be used; day $\langle 8/9/94 \rangle$ must then be checked.

(8)

(6) RA tree:



Plan:



Cost:

$$500 + 500 * 1000 / 10;$$

(7) No index:

